

ML5G-PS-005: Network failure prediction on CNFs 5GC with Linux eBPF

# **Regression-based Practical Network Failure Prediction on 5G Core Network Using AutoML**

Team MLAB-NFP

Member: Jiwon Lee, Bojian Du, Kentaro Matsuura, and Ryoma Kondo  
Morikawa Narusue Laboratory, the University of Tokyo

# Background

## Challenge

**How early** and accurately ( $f1 \text{ score} \geq 0.9$ ) the future network failures can be predicted using 3 types of metrics: Basic (cAdvisor), Fine-grained (eBPF), 5G metric

- Target value for prediction: the number of registration failure at 10 min

## Task

### Task 1


- **Use all (3326) metrics** for prediction

### Task 2

- **Select a part of metrics** for prediction

## Dataset

- 10 seconds interval x 70 rows per cycle
- Normal cycle (75%) / Abnormal cycle (25%)
- Failure event starts at 1.5 min (90 sec)



# Task1: Objective


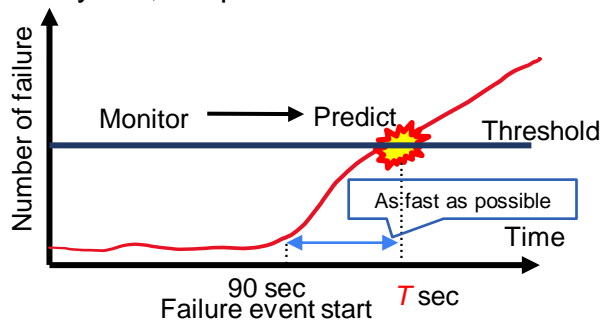
## Our Concept

In practice, the start time of failure events should be **unknown**

- It is **NOT appropriate** to predict the condition of cycles **only at a specific time  $T$** 
  - The start time of failure events can be later than  $T$
  - Can't determine the cycle label when model raise alarm after  $T$

We aim to provide a method which can apply failure prediction **for all data by time order up to 600sec** in every cycle

- For abnormal cycles, the prediction values **should exceed threshold as early as possible**
- For normal cycles, the prediction values **should not exceed threshold for all data up to 600sec**



# Task1: Methodology Summary

## Model

AutoGluon-Tabular (**regression model with stacking Ensemble**)

## Training Data Tuning

- Normal cycles:  
**Downsampling for the whole cycle**
- Abnormal cycles:  
**Using the data at early stage of failure only**

## Output Tuning

**Number of registration failures at 10min for the cycle corresponding to every input**

## Threshold Tuning

**Maximizing the margin between the prediction result of normal and abnormal cycles**

# Task1: Methodology Details

## Model

## AutoGluon-Tabular[1]

An open-source AutoML (Automated Machine Learning) framework

### Strengths

- **Ensemble** multiple models and **stack** them in multiple layer
- Can robustly take raw data and deliver **high-quality predictions without any user input**
  - Beat 99% of the participating data scientists in Kaggle competitions [1]

### Model training setting

- `problem_type='regression'`
- `presets='good_quality'`
- `time_limit=14400 [sec]`

### Auto selected submodels

- LightGBMLarge
- LightGBM
- LightGBMXT
- CatBoost
- XGBoost
- Random forest
- ExtraTrees




Fig. AutoGluon's multi-layer stacking strategy

[1] Erickson N, Mueller J, Shirkov A, et al. Autogluon-tabular: Robust and accurate automl for structured data[J]. arXiv preprint arXiv:2003.06505, 2020.

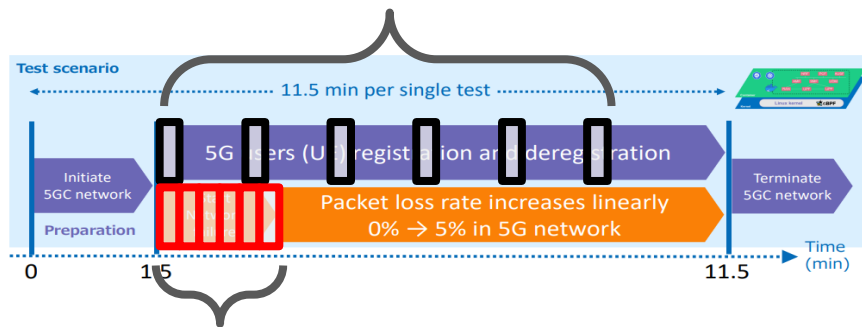
# Task1: Methodology Details

## Training Data Tuning

Different strategies for normal and abnormal cycles

### Normal cycles (450 cycles × 6 rows)

- **Downsampling for the whole cycle**
  - Constant interval of 100 sec
  - i.e. data with row index [0,10,20,30,40,50]



### Abnormal cycles (150 cycles × 6 rows)

- **Using the data at early stage of failure only**
  - i.e. data with row index [10,11,12,13,14,15]

- **Purpose**

- **Reduce false alarm after the prediction time** by learning the data from the whole cycle
- **Avoid data imbalance** in training phase

- **Purpose**

- **Achieve early prediction** by learn the indication of abnormal at the early stage of failure event only

# Task1: Methodology Details

## Output Tuning

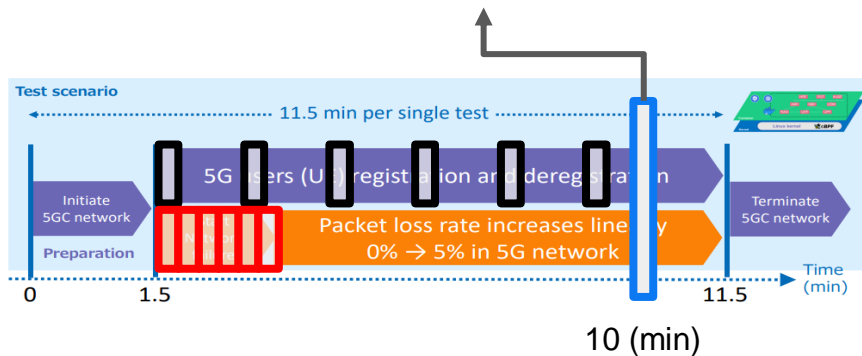
Number of registration failures at 10min for the cycle corresponding to every input

### Output data for each input

- **Number of registration failures at 10min,** instead of at current time, **of the same cycle as the input**
  - i.e., 6 same value for 6 inputs from the same cycle

- **Purpose**

- Let the model to correctly predict number of registration failures at 10min **at any time** in a cycle
  - Because the number of registration failures at the current time of input is not concerned in this task




# Task1: Methodology Details

## Threshold Tuning


Maximizing the margin between the prediction result of normal and abnormal cycles

- Based on the prediction of training dataset (600 cycles × 60 rows)



Histograms of **the predicted number of registration failures at 10min** in training cycles

a) Normal cycles



b) Abnormal cycles (after the start of failure events)

## Threshold

- The **mean value** of the 99th percentile of predicted failures in normal cycles (16.75) and the 1st percentile of predicted failures in abnormal cycles (34.57): **25.66**
  - The percentiles are used for excluding 1% outliers



# Task1: Result Summary

Prediction time

120sec

F1 score

At 120sec

Precision: **0.866**  
Recall: **0.947**

0.904

Up to 600sec

Precision: **0.834**  
Recall: **1.000**

0.904

- Achieve **100%** detection rate
- Keep a low false alarm rate of **7.1%**

## Confusion matrix

	Predicted Positive	Predicted Negative
Actual Positive	71	4
Actual Negative	11	214

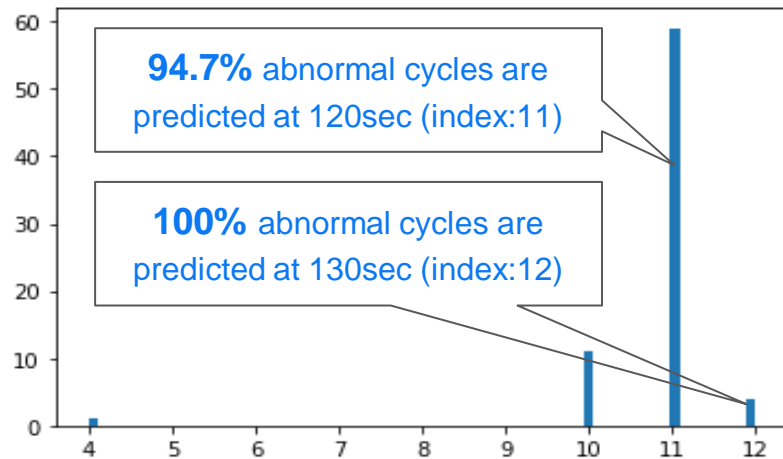

	Predicted Positive	Predicted Negative
Actual Positive	75	0
Actual Negative	16	209

# Task1: Result Details

Our model not only achieves a high F1 score (0.904) at the prediction time (120sec), but also keeps high score (0.904) after that time up to 600sec

Confusion matrix before / after 120sec

F1:0.904 / 0.904		Predicted	
		Positive	Negative
Actual	Positive	71 / 75	4 / 0
	Negative	11 / 16	214 / 209



Histograms show **the row index (0~59)** of **the first time the prediction value exceeds threshold** on test dataset

a) Normal cycles

b) Abnormal cycles

# Background

## Challenge

**How early** and accurately ( $f1 \text{ score} \geq 0.9$ ) the future network failures can be predicted using 3 types of metrics: Basic (cAdvisor), Fine-grained (eBPF), 5G metric

- The target value for prediction: the number of registration failure at 10 min

## Task

### Task 1


- Use all (3326) metrics for prediction

### Task 2

- Select a part of metrics for prediction

## Dataset

- 10 seconds interval x 70 rows per cycle
- Normal cycle (75%) / Abnormal cycle (25%)
- Failure event starts at 1.5 min (90 sec)



# Task2: Feature selection Details

## Random Forest Feature Importance of the Model in Task1

✖ Due to the limitation of computation resource, we calculated feature importance twice by randomly selecting 400 rows from training data (keeping normal:abnormal =3:1) for each time

✖ Features whose importance are lower than “**index**” are considered less important

- Features referring to **UDM-TCP RTT** and **cAdvisor-container memory RSS** tend to be important

→ **86 columns referring to UDM-TCP RTT and 10 columns referring to cAdvisor-container memory RSS are selected**

	importance
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.80.hist.bins._64_127.count	1.076865
smf.smf.app.cadvisor.container_memory_rss	0.931215
ausf.ausf.app.cadvisor.container_memory_rss	0.757510
amf.amf.app.cadvisor.container_memory_rss	0.749704
upf.upf3.app.cadvisor.container_memory_rss	0.593211
amf.amf.infra.tcpwin_192.168.13.80_192.168.13.82.snd_cwnd.hist.bins._4_7.count	0.559771
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.80.stat.avg	0.444564
udm.udm.app.cadvisor.container_memory_rss	0.405376
amf.amf.infra.tcpwin_192.168.13.80_192.168.13.82.snd_cwnd.hist.stat.min	0.351711
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.82.stat.avg	0.306957
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.80.stat.max	0.261798
upf.upf1.infra.runqlat.hist.bins._256_511.count	0.250903
amf.amf.infra.tcpwin_192.168.13.80_192.168.13.72.snd_cwnd.hist.stat.min	0.243320
nrf.nrf.app.cadvisor.container_memory_working_set_bytes	0.176476
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.82.stat.max	0.161863
<b>index</b>	0.149615

	importance
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.80.hist.bins._64_127.count	1.274811
smf.smf.app.cadvisor.container_memory_rss	1.253748
amf.amf.app.cadvisor.container_memory_rss	1.195922
ausf.ausf.app.cadvisor.container_memory_rss	1.081871
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.80.stat.avg	0.653811
udm.udm.app.cadvisor.container_memory_rss	0.536127
upf.upf3.app.cadvisor.container_memory_rss	0.515652
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.82.stat.avg	0.500319
amf.amf.infra.tcpwin_192.168.13.80_192.168.13.82.snd_cwnd.hist.bins._4_7.count	0.386906
upf.upf2.app.cadvisor.container_memory_rss	0.378289
amf.amf.infra.tcprrt_192.168.13.80_192.168.13.70.stat.avg	0.313949
amf.amf.infra.tcpwin_192.168.13.80_192.168.13.82.snd_cwnd.hist.stat.min	0.290665
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.82.stat.max	0.286260
upf.upf1.app.cadvisor.container_memory_rss	0.265179
igw.igw2.app.cadvisor.container_memory_working_set_bytes	0.250817
upf.upf1.infra.runqlat.hist.bins._256_511.count	0.244296
nrf.nrf.app.cadvisor.container_memory_working_set_bytes	0.224415
igw.igw2.infra.tcprrt_192.168.14.45_192.168.12.1.hist.bins._8192_16383.count	0.204181
udm.udm.infra.tcprrt_192.168.13.70_192.168.13.80.stat.max	0.192238
smf.smf.infra.tcprrt_192.168.13.82_192.168.13.70.stat.min	0.188466
amf.amf.infra.tcprrt_192.168.13.80_192.168.13.70.stat.max	0.178001
upf.upf1.app.cadvisor.devices._/dev/sdb2.container_fs_writes_merged	0.170604
upf.upf3.infra.runqlat.hist.bins._128_255.count	0.167914
<b>index</b>	0.167500

# Task2: Feature selection Details

## Thesis on Important Features

### Definition

UDM TCP RTT

- UDM : Unified Data Management
- RTT : Round-Trip Time
  - Time for a signal/data to be sent to the other end of the communication until a response is returned.

### Hypothesis

- UDM – subscriber management, user identification processing  
→ **related to registration**
- RTT – depends on physical distance, number of devices, processing time  
→ **Increase when failure occur**

### Definition

Container memory RSS


- RSS : Resident set size
  - Memory usage
  - Physical memory consumption

### Hypothesis

- **Failure event will consume more physical memory than usual**


# Task2: First Try

## Threshold Tuning




Histograms of the predicted number of registration failures at 10min in training cycles

a) Normal cycles



b) Abnormal cycles

(after the start of failure events, index:10~59)



c) Abnormal cycles

(only the early stage of failure events, index:10~15)

## Modification from Task1

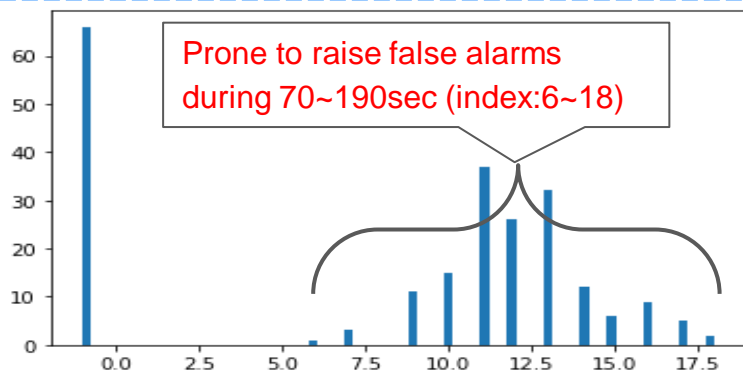
- **Use only the early stage of failure events for abnormal cycles**
  - As the prediction of the later stage of failure events are not accurate (too low)
    - It won't affect prediction performance, since the prediction values should exceed threshold at the early stage
- **Exclude 5% outliers instead of 1% outliers in the calculation of threshold**
  - As the distribution of the prediction values becomes more dispersed for both normal and abnormal cycles

# Task2: First Try

The new model **keeps a high prediction rate from 120sec**  
however, it **tends to raise more false alarms during 70~190sec**

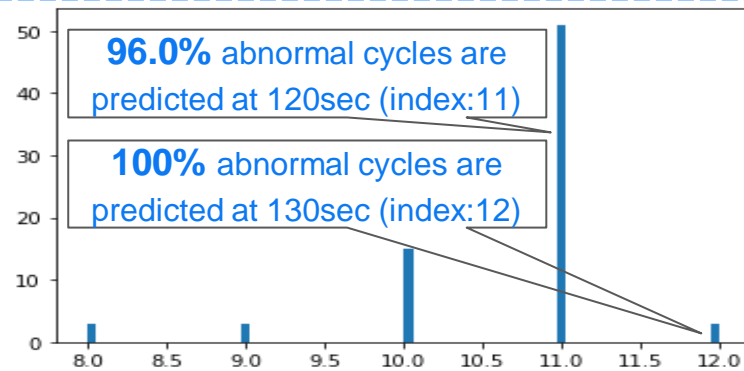
Confusion matrix before / after 120sec

F1:0.673 / 0.485		Predicted	
		Positive	Negative
Actual	Positive	72 / 75	3 / 0
	Negative	67 / 159	158 / 66



Histograms of the row index (0~59) of the first time the prediction value exceeds threshold on test dataset

a) Normal cycles



b) Abnormal cycles

## Discussion

- Features referring to **UDM-TCP RTT** and **cAdvisor-container memory RSS** shows the potential in predicting abnormal cycles
- False alarms in normal cycles are tend to occur in **the same time of the training periods of abnormal cycles**, as the beginning time of failure events are almost the same in this dataset
  - Adding more training samples** or **providing more variety in failure event** might help reduce this problem

# Task2: Methodology Details

## Modification of training data

- Separate training data (600 cycles) by 500:100 ratio and create validation data
  - Use training data (500 cycles) for training data tuning
    - Double the rows of training data
  - Use validation data (100 cycles) for threshold tuning
    - Find the value that maximize f1 score




# Task2: Methodology Details

## Training Data Tuning

Different strategies for normal and abnormal cycles

### Normal cycles (375 cycles × 12 rows)

- **Downsampling for the whole cycle**
  - Constant interval of 50 sec
  - i.e. data with row index [0,5,10,15,20,25,30,35,40,45,50,55]



- **Purpose**

- Improve prediction on normal cycles by adding more training samples

### Abnormal cycles (125 cycles × 12 rows)


- **Using the data at early stage of failure only**
  - i.e. data with row index [10,11,12,13,14,15,16,17,18,19,20,21]

# Task2: Methodology Details

## Threshold Tuning

Crossing point of the f1 score of 13<sup>th</sup> row and 60<sup>th</sup> row

- Based on the validation dataset of training data (100 cycles × 70 rows)



F1 scores of threshold from 20 to 60 at 13<sup>th</sup> and 60<sup>th</sup> row

- Change the threshold from 20 ~ 60 for both 13<sup>th</sup> row and 60<sup>th</sup> row
  - Check the point where the two graphs cross each other
  - 13<sup>th</sup> row: threshold = 40, f1 score = 0.912
  - 60<sup>th</sup> row: threshold = 40, f1 score = 0.943

## Threshold

- The **crossing point** of the f1 score of the 60<sup>th</sup> row (red) and the f1 score of the 13<sup>th</sup> row (blue): **40.00**
  - The crossing point used for setting threshold to predict with high accuracy throughout the whole cycle

# Task2: Result Summary

## Selected Features

~~3326~~ → 96 columns

Features referring to UDM-TCP RTT (86 columns)  
and cAdvisor-container memory RSS (10 columns)

## Prediction time

96.0% abnormal cycles  
are predicted at

130sec

## F1 score

### At 130sec

Precision: 0.907  
Recall: 0.907

0.907

### Up to 600sec

Precision: 0.915  
Recall: 1.000

0.955

## Confusion matrix

	Predicted Positive	Predicted Negative
Actual Positive	68	7
Actual Negative	7	218

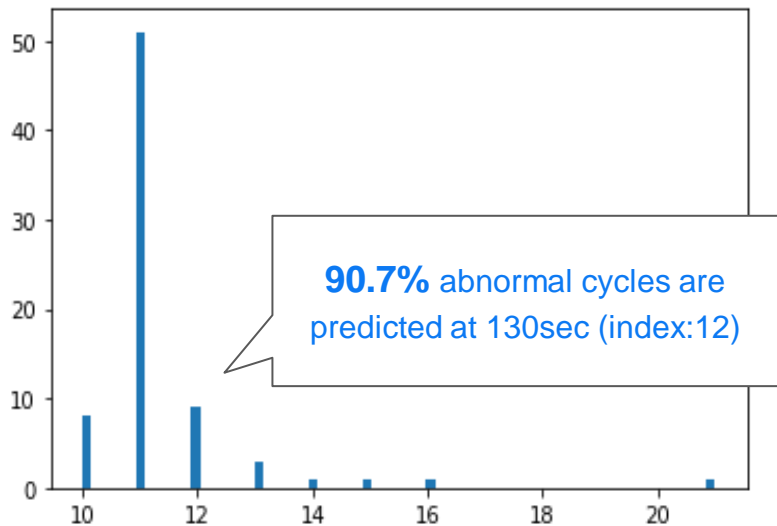

	Predicted Positive	Predicted Negative
Actual Positive	75	0
Actual Negative	7	218

# Task2: Result Details

Our model achieved a high F1 score both at the prediction time 130sec (0.907), and after that time up to 600sec (0.970).

Confusion matrix before / after 130sec

F1:0.907 / 0.955		Predicted	
		Positive	Negative
Actual	Positive	68 / 75	7 / 0
	Negative	7 / 7	218 / 218



Histograms of the row index (0~59) of the first time the prediction value exceeds threshold on test dataset

a) Normal cycles

b) Abnormal cycles

# Summary

## Task 1

### Method

- Apply prediction for all data by time up to 600 sec in every cycle
- Model: AutoGluon-Tabular
  - Training data: downsample to 6 rows per cycle
  - Output data: failures at 10min

### Result

- Prediction time: **120 sec**
- F1 score:
  - At 120 sec: **0.904**
  - Up to 600 sec: **0.904**

## Task 2

### Method

- Feature selection: UDM-TCP RTT (86 columns) and cAdvisor-container memory RSS (10 columns)
- Model: AutoGluon-Tabular
  - Training data: downsample to 12 rows per cycle
  - Output data: failures at 10min

### Result

- Prediction time: **130 sec**
- F1 score:
  - At 130 sec: **0.907**
  - Up to 600 sec: **0.955**